🍄 Machine Learning Project Report: Mushroom Edibility Classification

1. Project Goal and Overview

The main objective of this project was to determine whether a mushroom is **edible** or **poisonous** using data science and machine learning techniques. We started with a raw dataset of mushroom characteristics and followed a standard pipeline: **data cleaning**, **preprocessing**, **exploratory analysis**, **dimensionality reduction (PCA)**, **unsupervised learning (K-Means)**, and finally, **supervised classification (Random Forest)**.

---

2. Preparing the Data (The Cleanup Phase)

Before any machine learning can happen, the data must be clean and organized.

Competency	Action Taken	Why it Matters
Data Cleaning	**Missing values** (represented by '?') were filled in with the most common value (mode). A feature with only **one unique value** (the 'p.2' feature) was **removed** as it provided no information.	Ensures the data is complete and every feature contributes unique information.
Preprocessing	All categorical codes (e.g., 'k' for black, 'e' for edible) were converted into **numerical columns** using **One-Hot Encoding** and **Label Encoding**.	Machine learning models only understand numbers. This step expanded our initial 22 columns to **116 numerical features**.
Data Split	The data was split into a **Training Set (67%)** and a **Testing Set (33%)**.	The model trains on the Training Set and is tested on unseen data (the Testing Set) to ensure it performs well in the real world.

---

3. Exploring Data Structure and Complexity

With 116 features, the data is complex. We used **Principal Component Analysis (PCA)** to simplify and visualize it.

A. Dimensionality Reduction (PCA)

PCA helps us find the most important combinations of features that capture the most variation in the data.

- **Visualization:** We reduced the data to just **two components** to create a 2D scatter plot, colored by the target (edible/poisonous).

    - **Finding:** The plot showed large regions where the two mushroom types were perfectly separated, suggesting the classification task would be easy.
- **Information Retention:** We checked how many components are actually needed to keep
- 95%
- **95%** of the original information.
    - **Finding:** We still needed **109 out of the 116 components**. This means while the task is simple for a classifier, the features themselves are largely **non-redundant**. The complexity (116 features) is justified.

B. Unsupervised Clustering (K-Means)

We then used **K-Means clustering** (unsupervised learning) to see how the data naturally groups itself, without using the 'edible/poisonous' labels.

- **Optimal Clusters:** The **Elbow Method** confirmed that
- $K=2$
- **$K=2$** clusters were most appropriate (matching our two target classes).
- **Finding:** When plotting the clusters against the true labels, the K-Means groups were highly **mixed** (
- $\approx 50/50$
- $\approx 50/50$ split) .
- **Implication:** This shows that the factors that make a mushroom edible/poisonous **do not perfectly align** with the factors that make the mushroom features *look* similar. This is where a supervised model excels.

---

4. Supervised Classification and Final Results

We chose the **Random Forest** model—an Ensemble Method known for accuracy—to classify the mushrooms.

Competency	Model Used	Result
Classification Model	**Random Forest Classifier** (trained on standardized features)	**Test Accuracy:** 1.0000 **1.0000 (100%)**
Performance Check	Retrained Random Forest using the 109 **109** PCA components.	**Test Accuracy:** 1.0000 **1.0000 (100%)**

Conclusion

The **Random Forest** model achieved **perfect accuracy** (

100%

100%) in distinguishing between edible and poisonous mushrooms. This confirms that even though the dataset has many features, there is a **strong, clean signal** within the data that allows for flawless separation. The model is highly reliable for this classification task.

All figures and reports produced during this analysis are saved in the `results/figures/` and `results/reports/` directories for full reproducibility.